
Assignment 1 - Group 18

k-Nearest Neighbors for Alzheimer's Disease Classification

Isaac Edem Adoboe (202384695)
Department of Mathematics and Statistics
Memorial University of Newfoundland

ieadoboe@mun.ca

Blessing Ijeoma Benjamin-Igwe (202384108)
Department of Mathematics and Statistics
Memorial University of Newfoundland

bbenjaminigw@mun.ca

Promee Shankar Kundu (202382587)
Department of Computer Science
Memorial University of Newfoundland

pskundu@mun.ca

Introduction

The goal of this assignment was to examine how model capacity, controlled by the neighborhood size k , influences classification performance when distinguishing between **stable normal controls (sNC)** and **stable Alzheimer's disease patients (sDAT)** using glucose metabolism measurements from the **isthmuscingulate** and **precuneus** regions.

Question 1: Effect of Model Capacity on kNN Classification (Euclidean Distance)

k-nearest neighbor (kNN) classifiers were trained using the **Euclidean distance metric** for

$$k \in \{1, 3, 5, 10, 20, 30, 50, 100, 150, 200\}.$$

Decision boundaries were visualized on a two-dimensional grid for each value of k . Training samples are indicated with circular markers and test samples with plus markers. The sNC class is shown in green and the sDAT class in blue. Training and test error rates are reported directly in the figure titles.

Summary of Error Rates

To complement the visual analysis, representative training and test error rates are summarized below:

k	Train Error	Test Error
1	0.0000	0.2200
3	0.1435	0.2050
5	0.1561	0.1700
10	0.1519	0.1700
20	0.1688	0.1700
30	0.1688	0.1600
50	0.1582	0.1900
100	0.1962	0.2000
150	0.1920	0.1900
200	0.2215	0.2050

This table highlights the transition from overfitting to underfitting as k increases.

Overfitting Regime (Low Bias, High Variance)




Figure 1: Euclidean kNN decision boundary for $k = 1$. The highly irregular boundary closely follows individual training samples, resulting in near-zero training error but elevated test error, characteristic of overfitting.

For $k = 1$, the classifier exhibits **low bias and high variance**. The decision boundary adapts excessively to individual data points, capturing noise rather than generalizable structure. This leads to excellent training performance but poor generalization.

Optimal Regime (Balanced Bias–Variance Tradeoff)




Figure 2: Euclidean kNN decision boundary for $k = 30$. This value of k yields the lowest observed test error and produces a smooth yet flexible boundary that best balances bias and variance.

At intermediate values of k , the classifier achieves its best generalization performance. For $k = 30$, the decision boundary remains responsive to meaningful structure while avoiding excessive sensitivity to noise. This value of k corresponds to the **minimum test error** observed in the experiments and represents the optimal model capacity.

Underfitting Regime (High Bias, Low Variance)




Figure 3: Euclidean kNN decision boundary for $k = 200$. The overly smooth boundary fails to capture local structure, leading to increased training and test error due to underfitting.

For very large values of k , the classifier exhibits **high bias and low variance**. The decision boundary becomes overly simplistic, and both training and test errors increase and converge, indicating underfitting.

Overall, these results demonstrate that **model capacity decreases as k increases**, and that optimal performance is achieved at an intermediate value of k .

Question 2: Effect of Distance Metric

Using the optimal neighborhood size $k = 30$, a kNN classifier was trained with the **Manhattan (L1) distance metric**.




Figure 4: Manhattan kNN decision boundary for $k = 30$. The Manhattan-distance classifier produces more axis-aligned boundaries and achieves test performance comparable to, but slightly worse than, the Euclidean-distance classifier.

While both distance metrics show similar qualitative behavior, Euclidean distance produces smoother, radially symmetric decision boundaries, whereas Manhattan distance yields axis-aligned boundaries due to its reliance on absolute coordinate differences. Given its lower observed test error, **Euclidean distance** was selected for subsequent analysis. These results suggest that, for this dataset, the choice of k has a greater impact on performance than the choice of distance metric.

Distance Metric	Train Error	Test Error
Euclidean ($k=30$)	0.1688	0.1600
Manhattan ($k=30$)	0.1646	0.1650

Question 3: Error Rate Versus Model Capacity

Using Euclidean distance, model performance was analyzed as a function of **model capacity**, defined as $\frac{1}{k}$. Values of k were chosen to span the range $\frac{1}{k} \in [0.01, 1.0]$. Training and test error rates were plotted against model capacity on a logarithmic x-axis.




Figure 5: Training and test error rates versus model capacity ($1/k$) using Euclidean distance.

The test error curve exhibits a U-shaped profile, illustrating the bias-variance tradeoff. At high capacity (small k), test error is elevated due to overfitting. As capacity decreases, test error reaches a minimum at $k = 30$ (**capacity** ≈ 0.033) achieving a test error of **0.1600**. Further reductions in capacity result in underfitting and increased error due to high bias.

The Bayes classifier error is not shown, as it cannot be estimated without knowledge of the true data-generating distribution.

Question 4: Design of the Best kNN Classifier

Based on the experimental results obtained in Questions 1–3, a final k-nearest neighbors (kNN) classifier was designed to optimize generalization performance for distinguishing between stable normal controls (sNC) and stable Alzheimer’s disease patients (sDAT).

From the bias–variance analysis conducted in Question 1 and the capacity analysis in Question 3, a neighborhood size of $k = 30$ was selected, as it consistently achieved the lowest test error while maintaining a stable and interpretable decision boundary. This value of k provides an effective balance between overfitting at small k and underfitting at large k .

The final classifier employs the **Euclidean distance metric**, which demonstrated slightly better performance than the Manhattan distance metric in Question 2 and produced smoother, radially symmetric decision boundaries. To further improve classification robustness beyond the experiments in Questions 1–3, **feature standardization** was applied using z-score normalization. This ensures that glucose metabolism measurements from the isthmus cingulate and precuneus regions contribute equally to distance computations and prevents dominance by differences in scale.

In addition, **distance-weighted voting** was used so that closer neighbors exert greater influence on classification decisions. This refinement improves local sensitivity while preserving the overall bias–variance balance achieved with $k = 30$.

The classifier was implemented in the function:

```
diagnoseDAT(Xtest, data_dir)
```

To maximize generalization performance on the independent test dataset, the final classifier was trained on **all available labeled data** (combining the 474 training samples and 200 test samples provided for Questions 1–3), yielding a total training set of 674 samples. This choice to include the test set was made because the `diagnoseDAT()` function would be evaluated on an independent test dataset not available during development. This approach follows standard machine learning practice: once hyperparameters have been selected via proper train-test validation, the model is retrained on all available data to improve robustness for deployment.

This function loads the provided labeled datasets, constructs a standardized training set, trains the optimized kNN model, and returns predicted class labels (0 for sNC and 1 for sDAT) for unseen test samples. The implementation is fully self-contained and designed to generalize well to the independent blinded test dataset used for grading.

Overall, the final classifier reflects a principled design informed by systematic experimentation, balancing model complexity, distance metric choice, and preprocessing to achieve strong generalization performance.

Conclusion

This assignment demonstrates how kNN classification performance is governed by the interaction between model capacity, distance metric, and preprocessing choices. Through systematic experimentation with neighborhood sizes ranging from $k=1$ to $k=200$, we observed the classic bias-variance tradeoff: small k values led to overfitting (0% training error, 22% test error at $k=1$), while large k values caused underfitting (22% training error, 20.5% test error at $k=200$).

The optimal configuration achieved **16% test error at $k=30$** using Euclidean distance with feature standardization and distance-weighted voting. Comparison with Manhattan distance (Q2) revealed minimal performance differences, suggesting that model capacity (choice of k) has greater impact than distance metric selection for this dataset.

The final optimized classifier, trained on all 674 available samples and submitted for evaluation on an independent test dataset, demonstrates that principled hyperparameter selection through systematic experimentation yields robust and generalizable models for medical diagnosis applications.

Figure	k value / Metric	Description
Figure 1	k=1, Euclidean	Overfitting regime
Figure 2	k=30, Euclidean	Optimal regime
Figure 3	k=200, Euclidean	Underfitting regime
Figure 4	k=30, Manhattan	Distance metric comparison
Figure 5	All k values	Bias-variance analysis

Appendix: Decision Boundary Visualizations

This appendix presents the complete set of decision boundary visualizations for all values of k tested in Question 1, using the Euclidean distance metric.

Small k values ($k = 1, 3, 5, 10$)






Figure 6: Decision boundaries for small k values showing overfitting regime (high variance, low bias).


Moderate k values ($k = 20, 30, 50, 100$)




(a) $k = 20$



(b) $k = 30$ (optimal)




(c) $k = 50$




(d) $k = 100$

Figure 7: Decision boundaries for moderate k values showing optimal regime and transition to underfitting.

Large k values ($k = 150, 200$)



(a) $k = 150$



(b) $k = 200$

Figure 8: Decision boundaries for large k values showing underfitting regime (high bias, low variance).